

loan-data-prediction-1

January 14, 2024

[]:

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
from IPython import display
```

Loading the data

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Copy the path of the file in the drive. You can also change the Directory.

```
[4]: df = pd.read_csv('/content/drive/MyDrive/datasets/credit_risk_dataset.csv')
```

```
[5]: df.head(10)
```

```
[5]:   person_age  person_income  person_home_ownership  person_emp_length  \
0          22         59000             RENT             123.0
1          21          9600              OWN              5.0
2          25          9600         MORTGAGE              1.0
3          23         65500             RENT              4.0
4          24         54400             RENT              8.0
5          21          9900              OWN              2.0
6          26         77100             RENT              8.0
7          24         78956             RENT              5.0
8          24         83000             RENT              8.0
9          21         10000              OWN              6.0

   loan_intent  loan_grade  loan_amnt  loan_int_rate  loan_status  \
0  PERSONAL      D      35000      16.02            1
1  EDUCATION      B       1000      11.14            0
2  MEDICAL       C       5500      12.87            1
```

3	MEDICAL	C	35000	15.23	1
4	MEDICAL	C	35000	14.27	1
5	VENTURE	A	2500	7.14	1
6	EDUCATION	B	35000	12.42	1
7	MEDICAL	B	35000	11.11	1
8	PERSONAL	A	35000	8.90	1
9	VENTURE	D	1600	14.74	1

	loan_percent_income	cb_person_default_on_file	cb_person_cred_hist_length
0	0.59	Y	3
1	0.10	N	2
2	0.57	N	3
3	0.53	N	2
4	0.55	Y	4
5	0.25	N	2
6	0.45	N	3
7	0.44	N	4
8	0.42	N	2
9	0.16	N	3

Size and information of the dataset

```
[6]: df.shape
```

```
[6]: (32581, 12)
```

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32581 entries, 0 to 32580
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   person_age                            32581 non-null  int64
1   person_income                         32581 non-null  int64
2   person_home_ownership                 32581 non-null  object
3   person_emp_length                     31686 non-null  float64
4   loan_intent                           32581 non-null  object
5   loan_grade                           32581 non-null  object
6   loan_amnt                             32581 non-null  int64
7   loan_int_rate                         29465 non-null  float64
8   loan_status                           32581 non-null  int64
9   loan_percent_income                   32581 non-null  float64
10  cb_person_default_on_file              32581 non-null  object
11  cb_person_cred_hist_length             32581 non-null  int64
dtypes: float64(3), int64(5), object(4)
memory usage: 3.0+ MB
```

checking the null values in the dataset

```
[8]: df.isnull().sum()
```

```
[8]: person_age           0
     person_income       0
     person_home_ownership 0
     person_emp_length    895
     loan_intent          0
     loan_grade          0
     loan_amnt           0
     loan_int_rate       3116
     loan_status         0
     loan_percent_income  0
     cb_person_default_on_file 0
     cb_person_cred_hist_length 0
     dtype: int64
```

Filling null values with mode or median rather than deleting it due to more number of null values almost 10%.

```
[9]: print('person_emp_length mode {}'.format(df['person_emp_length'].mode()[0]))
     print('person_emp_length median {}'.format(df['person_emp_length'].median()))
     print('loan_int_rate mode {}'.format(df['loan_int_rate'].mode()[0]))
     print('loan_int_rate median {}'.format(df['loan_int_rate'].median()))
```

```
person_emp_length mode 0.0
person_emp_length median 4.0
loan_int_rate mode 7.51
loan_int_rate median 10.99
```

```
[10]: df['person_emp_length'].fillna(df['person_emp_length'].mode()[0], inplace=True)
     df['loan_int_rate'].fillna(df['loan_int_rate'].median(), inplace=True)
```

After filling the null values

```
[22]: df.isnull().sum()
```

```
[22]: person_age           0
     person_income       0
     person_home_ownership 0
     person_emp_length    0
     loan_intent          0
     loan_grade          0
     loan_amnt           0
     loan_int_rate       0
     loan_status         0
     loan_percent_income  0
```

```
cb_person_default_on_file    0
cb_person_cred_hist_length   0
dtype: int64
```

```
[11]: print('person_emp_length mode {}'.format(df['person_emp_length'].mode()[0]))
      print('person_emp_length median {}'.format(df['person_emp_length'].median()))
      print('loan_int_rate mode {}'.format(df['loan_int_rate'].mode()[0]))
      print('loan_int_rate median {}'.format(df['loan_int_rate'].median()))
```

```
person_emp_length mode 0.0
person_emp_length median 4.0
loan_int_rate mode 10.99
loan_int_rate median 10.99
```

There is change in mode of Loan interest rate after filling null values with median

```
[12]: df.dtypes #checking the datatype of values in the column
```

```
[12]: person_age                int64
      person_income            int64
      person_home_ownership    object
      person_emp_length        float64
      loan_intent              object
      loan_grade              object
      loan_amnt                int64
      loan_int_rate            float64
      loan_status              int64
      loan_percent_income       float64
      cb_person_default_on_file object
      cb_person_cred_hist_length int64
      dtype: object
```

Number of defaulter and non defaulter are to be checked to know the dataset

```
[13]: df['loan_status'].value_counts()
```

```
[13]: 0    25473
      1     7108
      Name: loan_status, dtype: int64
```

```
[14]: df['person_home_ownership'].unique()
```

```
[14]: array(['RENT', 'OWN', 'MORTGAGE', 'OTHER'], dtype=object)
```

```
[15]: df['loan_intent'].unique()
```

```
[15]: array(['PERSONAL', 'EDUCATION', 'MEDICAL', 'VENTURE', 'HOMEIMPROVEMENT',
          'DEBTCONSOLIDATION'], dtype=object)
```

```
[16]: df.describe()
```

```
[16]:
```


	person_age	person_income	person_emp_length	loan_amnt	\
count	32581.000000	3.258100e+04	32581.000000	32581.000000	
mean	27.734600	6.607485e+04	4.658114	9589.371106	
std	6.348078	6.198312e+04	4.159669	6322.086646	
min	20.000000	4.000000e+03	0.000000	500.000000	
25%	23.000000	3.850000e+04	2.000000	5000.000000	
50%	26.000000	5.500000e+04	4.000000	8000.000000	
75%	30.000000	7.920000e+04	7.000000	12200.000000	
max	144.000000	6.000000e+06	123.000000	35000.000000	

	loan_int_rate	loan_status	loan_percent_income	\
count	32581.000000	32581.000000	32581.000000	
mean	11.009620	0.218164	0.170203	
std	3.081611	0.413006	0.106782	
min	5.420000	0.000000	0.000000	
25%	8.490000	0.000000	0.090000	
50%	10.990000	0.000000	0.150000	
75%	13.110000	0.000000	0.230000	
max	23.220000	1.000000	0.830000	


	cb_person_cred_hist_length
count	32581.000000
mean	5.804211
std	4.055001
min	2.000000
25%	3.000000
50%	4.000000
75%	8.000000
max	30.000000

After observing the above values, there are some outlier which are impossible, i.e, MAX of persons age is 144 years, and MAX of person employment leneth is 123 years which cannot be possible.


```
[17]: plt.hist(df['loan_amnt'], bins=20)
plt.xlabel('Loan Amount')
plt.ylabel('Frequency')
plt.title('Distribution of Loan Amounts')
plt.show()
```



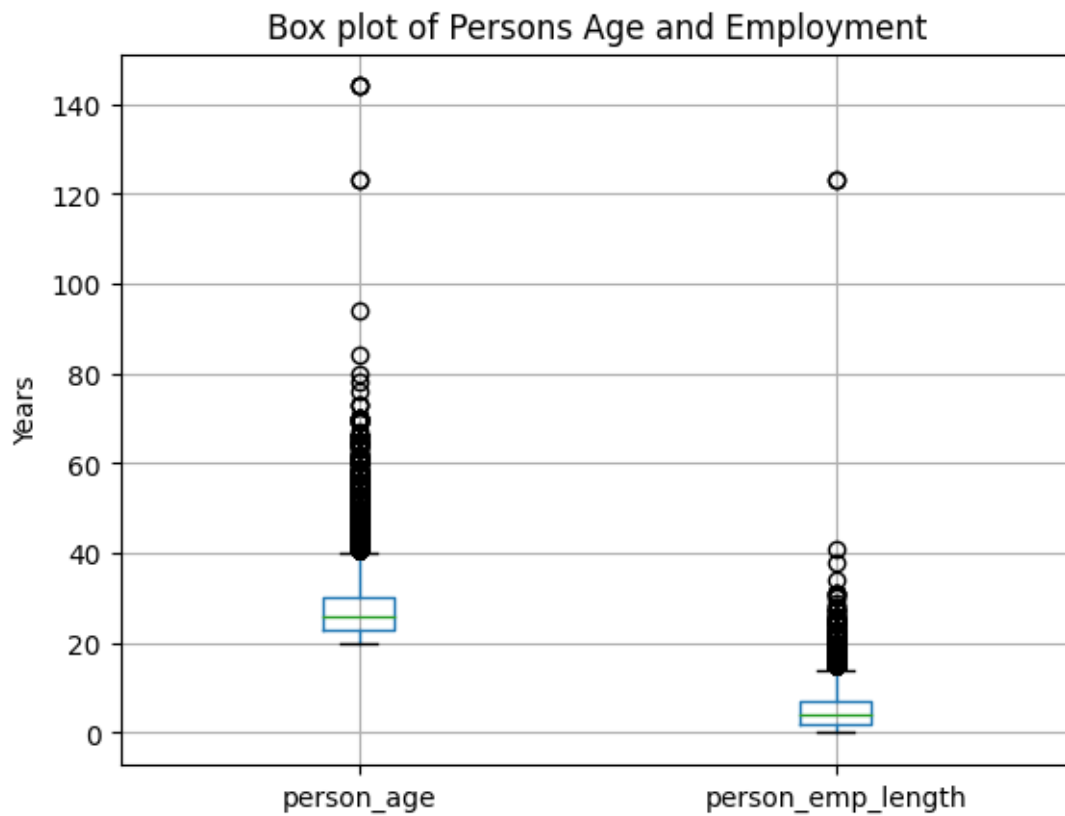
```
[18]: plt.hist(df['person_age'], bins=40)
plt.xlabel('AGE')
plt.ylabel('Frequency')
plt.title('Distribution of AGE')
plt.show()
```



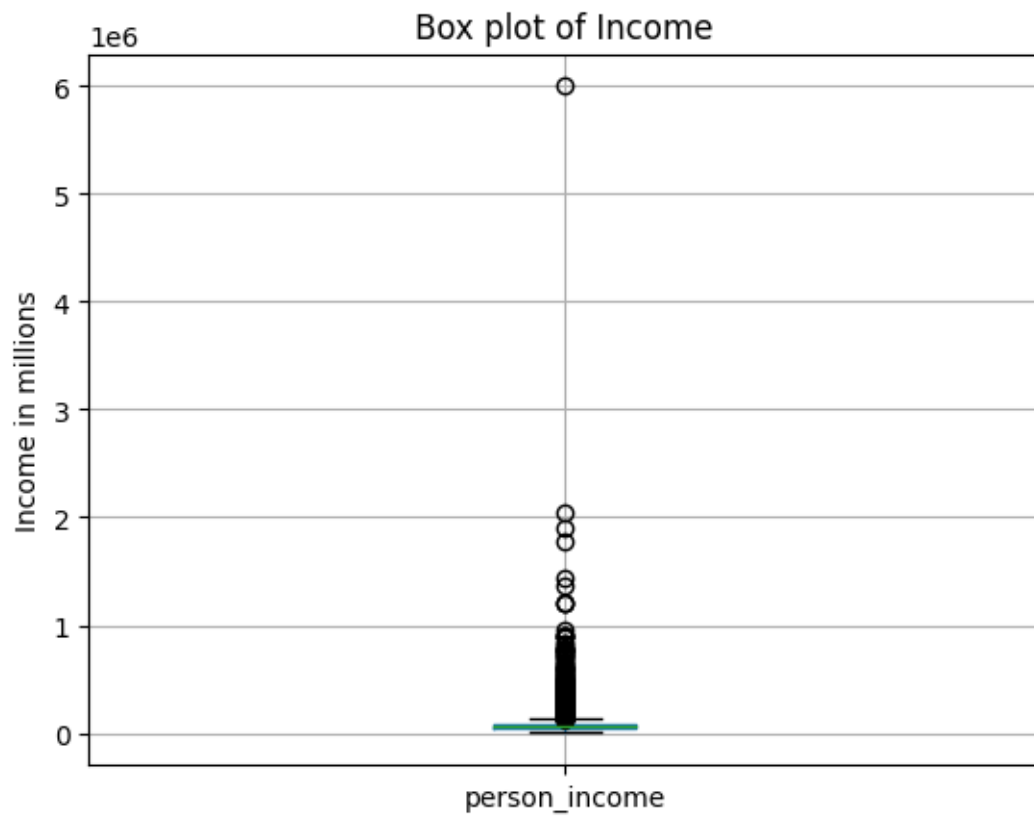
```
[19]: plt.hist(df['person_income'], bins=50)
plt.xlabel('Income')
plt.ylabel('Frequency')
plt.title('Distribution of Income')
plt.show()
```




```
[20]: df.boxplot(['person_age', 'person_emp_length'])  
plt.title('Box plot of Persons Age and Employment')  
plt.ylabel('Years')  
plt.show()
```

```
[21]: df.boxplot(['person_income'])  
plt.title('Box plot of Income')  
plt.ylabel('Income in millions')  
plt.show()
```



```
[22]: df.boxplot(['cb_person_cred_hist_length'])  
plt.show()
```



From above box plot we can observe the outliers

```
[23]: df = df[df['person_age']<=100]
      df = df[df['person_emp_length']<=60]
      df = df[df['person_income']<=3.000000e+06]
```

Removing the outliers using above condition

```
[24]: df.shape
```

```
[24]: (32574, 12)
```

Checking the correlation chart using the Heatmap.

```
[26]: f, ax = plt.subplots(figsize=(8,6))
      print(f)
      sns.heatmap(df.corr(),annot=True,linewidths=8,center =0,ax=ax,cmap="coolwarm")
```

Figure(800x600)

<ipython-input-26-1110c77e98d2>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only

to silence this warning.

```
sns.heatmap(df.corr(),annot=True,linewidths=8,center =0,ax=ax,cmap="coolwarm")
```

[26]: <Axes: >



One hot encoding of categorical columns.

```
[27]: dfohc= pd.get_dummies(df, columns=['person_home_ownership', 'loan_intent', 'loan_grade', 'cb_person_default_on_file'], drop_first=True)
dfohc.head(30)
```

```
[27]:
```

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	\
1	21	9600	5.0	1000	11.14	
2	25	9600	1.0	5500	12.87	
3	23	65500	4.0	35000	15.23	
4	24	54400	8.0	35000	14.27	
5	21	9900	2.0	2500	7.14	
6	26	77100	8.0	35000	12.42	

7	24	78956	5.0	35000	11.11
8	24	83000	8.0	35000	8.90
9	21	10000	6.0	1600	14.74
10	22	85000	6.0	35000	10.37
11	21	10000	2.0	4500	8.63
12	23	95000	2.0	35000	7.90
13	26	108160	4.0	35000	18.39
14	23	115000	2.0	35000	7.90
15	23	500000	7.0	30000	10.65
16	23	120000	0.0	35000	7.90
17	23	92111	7.0	35000	20.25
18	23	113000	8.0	35000	18.25
19	24	10800	8.0	1750	10.99
20	25	162500	2.0	35000	7.49
21	25	137000	9.0	34800	16.77
22	22	65000	4.0	34000	17.58
23	24	10980	0.0	1500	7.29
24	22	80000	3.0	33950	14.54
25	24	67746	8.0	33000	12.68
26	21	11000	3.0	4575	17.74
27	23	11000	0.0	1400	9.32
28	24	65000	6.0	32500	9.99
29	21	11389	5.0	4000	12.84
30	21	11520	5.0	2000	11.12

	loan_status	loan_percent_income	cb_person_cred_hist_length	\
1	0	0.10		2
2	1	0.57		3
3	1	0.53		2
4	1	0.55		4
5	1	0.25		2
6	1	0.45		3
7	1	0.44		4
8	1	0.42		2
9	1	0.16		3
10	1	0.41		4
11	1	0.45		2
12	1	0.37		2
13	1	0.32		4
14	0	0.30		4
15	0	0.06		3
16	0	0.29		4
17	1	0.32		4
18	1	0.31		4
19	1	0.16		2
20	0	0.22		4
21	0	0.25		2

22	1	0.52	4
23	0	0.14	3
24	1	0.42	4
25	1	0.49	3
26	1	0.42	3
27	0	0.13	3
28	1	0.50	3
29	1	0.35	2
30	1	0.17	3

	person_home_ownership_OTHER	person_home_ownership_OWN	...	\
1	0	1	...	
2	0	0	...	
3	0	0	...	
4	0	0	...	
5	0	1	...	
6	0	0	...	
7	0	0	...	
8	0	0	...	
9	0	1	...	
10	0	0	...	
11	0	1	...	
12	0	0	...	
13	0	0	...	
14	0	0	...	
15	0	0	...	
16	0	0	...	
17	0	0	...	
18	0	0	...	
19	0	0	...	
20	0	0	...	
21	0	0	...	
22	0	0	...	
23	0	1	...	
24	0	0	...	
25	0	0	...	
26	0	0	...	
27	0	1	...	
28	0	0	...	
29	1	0	...	
30	0	1	...	

	loan_intent_MEDICAL	loan_intent_PERSONAL	loan_intent_VENTURE	\
1	0	0	0	
2	1	0	0	
3	1	0	0	
4	1	0	0	

5	0	0	1
6	0	0	0
7	1	0	0
8	0	1	0
9	0	0	1
10	0	0	1
11	0	0	0
12	0	0	1
13	0	0	0
14	0	0	0
15	0	0	0
16	0	0	0
17	1	0	0
18	0	0	0
19	0	0	0
20	0	0	1
21	0	1	0
22	0	0	0
23	0	1	0
24	0	1	0
25	0	0	0
26	0	0	1
27	0	1	0
28	0	0	0
29	0	0	0
30	1	0	0

	loan_grade_B	loan_grade_C	loan_grade_D	loan_grade_E	loan_grade_F	\
1	1	0	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	
5	0	0	0	0	0	
6	1	0	0	0	0	
7	1	0	0	0	0	
8	0	0	0	0	0	
9	0	0	1	0	0	
10	1	0	0	0	0	
11	0	0	0	0	0	
12	0	0	0	0	0	
13	0	0	0	1	0	
14	0	0	0	0	0	
15	1	0	0	0	0	
16	0	0	0	0	0	
17	0	0	0	0	1	
18	0	0	1	0	0	
19	1	0	0	0	0	

20	0	0	0	0	0
21	0	0	0	1	0
22	0	0	1	0	0
23	0	0	0	0	0
24	0	0	1	0	0
25	0	1	0	0	0
26	0	0	0	1	0
27	0	0	0	0	0
28	1	0	0	0	0
29	0	1	0	0	0
30	1	0	0	0	0


	loan_grade_G	cb_person_default_on_file_Y
1	0	0
2	0	0
3	0	0
4	0	1
5	0	0
6	0	0
7	0	0
8	0	0
9	0	0
10	0	0
11	0	0
12	0	0
13	0	0
14	0	0
15	0	0
16	0	0
17	0	0
18	0	0
19	0	0
20	0	0
21	0	1
22	0	0
23	0	0
24	0	1
25	0	0
26	0	1
27	0	0
28	0	0
29	0	1
30	0	0

[30 rows x 23 columns]


```
[28]: f , ax = plt.subplots(figsize=(30,16))
print(f)
sns.heatmap(dfohc.corr(), annot=True,linewidths=8,center_
↪=0,ax=ax,cmap="coolwarm")
```

Figure(3000x1600)

[28]: <Axes: >



Separating the input and out (X & Y).

```
[29]: Y=dfohc['loan_status']
X= dfohc.drop('loan_status',axis=1)
```

[52]: X

```
[52]:      person_age  person_income  person_emp_length  loan_amnt  loan_int_rate \
1              21             9600                5.0        1000        11.14
2              25             9600                1.0         5500        12.87
3              23            65500                4.0       35000        15.23
4              24            54400                8.0       35000        14.27
5              21             9900                2.0        2500         7.14
...
32576          57            53000                1.0         5800        13.16
32577          54           120000                4.0       17625         7.49
```

32578	65	76000	3.0	35000	10.99
32579	56	150000	5.0	15000	11.48
32580	66	42000	2.0	6475	9.99

	loan_percent_income	cb_person_cred_hist_length	\
1	0.10	2	
2	0.57	3	
3	0.53	2	
4	0.55	4	
5	0.25	2	
...	
32576	0.11	30	
32577	0.15	19	
32578	0.46	28	
32579	0.10	26	
32580	0.15	30	

	person_home_ownership_OTHER	person_home_ownership_OWN	\
1	0	1	
2	0	0	
3	0	0	
4	0	0	
5	0	1	
...	
32576	0	0	
32577	0	0	
32578	0	0	
32579	0	0	
32580	0	0	

	person_home_ownership_RENT	...	loan_intent_MEDICAL	\
1	0	...	0	
2	0	...	1	
3	1	...	1	
4	1	...	1	
5	0	...	0	
...	
32576	0	...	0	
32577	0	...	0	
32578	1	...	0	
32579	0	...	0	
32580	1	...	1	

	loan_intent_PERSONAL	loan_intent_VENTURE	loan_grade_B	loan_grade_C	\
1	0	0	1	0	
2	0	0	0	1	
3	0	0	0	1	

4	0	0	0	1
5	0	1	0	0
...
32576	1	0	0	1
32577	1	0	0	0
32578	0	0	1	0
32579	1	0	1	0
32580	0	0	1	0

	loan_grade_D	loan_grade_E	loan_grade_F	loan_grade_G	\
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
5	0	0	0	0	
...	
32576	0	0	0	0	
32577	0	0	0	0	
32578	0	0	0	0	
32579	0	0	0	0	
32580	0	0	0	0	

	cb_person_default_on_file_Y
1	0
2	0
3	0
4	1
5	0
...	...
32576	0
32577	0
32578	0
32579	0
32580	0

[32574 rows x 22 columns]

Performing the Train test split of (0.75:0.25)

```
[30]: from sklearn.model_selection import train_test_split

X_train , X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.
↪25,random_state=10)
print(X_train.shape,Y_train.shape)
print(X_test.shape,Y_test.shape)
```

```
(24430, 22) (24430,)
(8144, 22) (8144,)
```

Data Transformation

```
[33]: from sklearn.preprocessing import StandardScaler
scale = StandardScaler()

X_train_trans = scale.fit_transform(X_train)
print(X_train_trans.shape)
X_test_trans = scale.transform(X_test)
print(X_test_trans.shape)
```

(24430, 22)

(8144, 22)

```
[34]: X_train_trans
```

```
[34]: array([[ -0.76047974, -0.26579741,  0.57866681, ..., -0.08542872,
        -0.04390413, -0.46178989],
       [ -0.76047974,  0.99037087,  0.57866681, ..., -0.08542872,
        -0.04390413, -0.46178989],
       [  0.37467066,  0.01442474,  0.08617155, ..., -0.08542872,
        -0.04390413, -0.46178989],
       ...,
       [  2.4828071 ,  2.20788783,  0.08617155, ..., -0.08542872,
        -0.04390413, -0.46178989],
       [-0.92264408, -0.69096206, -0.16007608, ..., -0.08542872,
        -0.04390413,  2.165487  ],
       [-0.76047974, -0.30444874,  0.08617155, ..., -0.08542872,
        -0.04390413, -0.46178989]])
```

The above code is creating array as output, To create dataframe below code is written

```
[35]: cols=['person_age', 'person_income', 'person_emp_length', 'loan_amnt', 'loan_int_rate', 'loan_perce
```

```
[36]: scale = StandardScaler()

X_train_transformed = pd.DataFrame(scale.
    ↪fit_transform(X_train), columns=cols, index = X_train.index)
X_test_transformed = pd.DataFrame(scale.transform(X_test), columns=cols, index =
    ↪X_test.index)

X_train_transformed.head()
```

```
[36]:
```

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	\
15001	-0.760480	-0.265797	0.578667	0.857754	1.055419	
14857	-0.760480	0.990371	0.578667	-0.931503	-0.087376	
23863	0.374671	0.014425	0.086172	-0.250636	-1.015897	
30338	1.509821	-0.265797	-0.898819	0.857754	-0.006212	
1884	-0.598315	-0.207820	-0.160076	0.841920	0.266501	

	loan_percent_income	cb_person_cred_hist_length	\
15001	1.119782	-0.692179	
14857	-1.313617	-0.692179	
23863	-0.471287	0.794343	
30338	1.119782	2.033111	
1884	0.932597	-0.692179	

	person_home_ownership_OTHER	person_home_ownership_OWN	\
15001	-0.056595	-0.294195	
14857	-0.056595	-0.294195	
23863	-0.056595	-0.294195	
30338	-0.056595	-0.294195	
1884	-0.056595	-0.294195	

	person_home_ownership_RENT	...	loan_intent_MEDICAL	\
15001	0.991522	...	-0.477632	
14857	-1.008551	...	2.093664	
23863	-1.008551	...	-0.477632	
30338	-1.008551	...	-0.477632	
1884	0.991522	...	-0.477632	

	loan_intent_PERSONAL	loan_intent_VENTURE	loan_grade_B	loan_grade_C	\
15001	-0.452432	2.183088	-0.691511	2.031759	
14857	-0.452432	-0.458067	1.446109	-0.492184	
23863	-0.452432	-0.458067	-0.691511	-0.492184	
30338	-0.452432	2.183088	-0.691511	-0.492184	
1884	2.210278	-0.458067	1.446109	-0.492184	

	loan_grade_D	loan_grade_E	loan_grade_F	loan_grade_G	\
15001	-0.356226	-0.173386	-0.085429	-0.043904	
14857	-0.356226	-0.173386	-0.085429	-0.043904	
23863	-0.356226	-0.173386	-0.085429	-0.043904	
30338	-0.356226	-0.173386	-0.085429	-0.043904	
1884	-0.356226	-0.173386	-0.085429	-0.043904	

	cb_person_default_on_file_Y
15001	-0.46179
14857	-0.46179
23863	-0.46179
30338	-0.46179
1884	-0.46179

[5 rows x 22 columns]

```
[37]: X_train_transformed.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 24430 entries, 15001 to 17679
Data columns (total 22 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   person_age                            24430 non-null  float64
 1   person_income                         24430 non-null  float64
 2   person_emp_length                     24430 non-null  float64
 3   loan_amnt                             24430 non-null  float64
 4   loan_int_rate                         24430 non-null  float64
 5   loan_percent_income                  24430 non-null  float64
 6   cb_person_cred_hist_length           24430 non-null  float64
 7   person_home_ownership_OTHER          24430 non-null  float64
 8   person_home_ownership_OWNS           24430 non-null  float64
 9   person_home_ownership_RENT           24430 non-null  float64
10   loan_intent_EDUCATION                 24430 non-null  float64
11   loan_intent_HOMEIMPROVEMENT           24430 non-null  float64
12   loan_intent_MEDICAL                   24430 non-null  float64
13   loan_intent_PERSONAL                  24430 non-null  float64
14   loan_intent_VENTURE                   24430 non-null  float64
15   loan_grade_B                         24430 non-null  float64
16   loan_grade_C                         24430 non-null  float64
17   loan_grade_D                         24430 non-null  float64
18   loan_grade_E                         24430 non-null  float64
19   loan_grade_F                         24430 non-null  float64
20   loan_grade_G                         24430 non-null  float64
21   cb_person_default_on_file_Y           24430 non-null  float64
dtypes: float64(22)
memory usage: 4.3 MB

```

Modeling

Logistic Regression

```
[38]: from sklearn.linear_model import LogisticRegression
```

```

LR_model = LogisticRegression()

LR_model.fit(X_train_transformed,Y_train)

y_pred_LR = LR_model.predict(X_test_transformed)

```

```
[41]: from sklearn import metrics

from sklearn.metrics import auc, accuracy_score, confusion_matrix,
     ↪ roc_auc_score, classification_report
```

```
print(classification_report(Y_test, LR_model.predict(X_test_transformed)))
LR_AUC=metrics.accuracy_score(Y_test,y_pred_LR)
print(LR_AUC)
```

	precision	recall	f1-score	support
0	0.89	0.95	0.92	6423
1	0.76	0.56	0.64	1721
accuracy			0.87	8144
macro avg	0.83	0.76	0.78	8144
weighted avg	0.86	0.87	0.86	8144

0.8699656188605108

KNN

```
[42]: from sklearn.neighbors import KNeighborsClassifier

KNNcls = KNeighborsClassifier()
KNNcls.fit(X_train_transformed,Y_train)
y_pred_KNN = KNNcls.predict(X_test_transformed)

print(classification_report(Y_test, KNNcls.predict(X_test_transformed)))
KNN_AUC=metrics.accuracy_score(Y_test,y_pred_KNN)
print(KNN_AUC)
```

	precision	recall	f1-score	support
0	0.91	0.97	0.94	6423
1	0.85	0.64	0.73	1721
accuracy			0.90	8144
macro avg	0.88	0.80	0.83	8144
weighted avg	0.90	0.90	0.89	8144

0.900171905697446

Support vector machines (SVM)

```
[43]: from sklearn.svm import SVC

SVMa = SVC()
SVMa.fit(X_train_transformed,Y_train)
y_pred_SVC = SVMa.predict(X_test_transformed)

print(classification_report(Y_test, SVMa.predict(X_test_transformed)))
SVM_AUC=metrics.accuracy_score(Y_test,y_pred_SVC)
print(SVM_AUC)
```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	6423
1	0.93	0.66	0.77	1721
accuracy			0.92	8144
macro avg	0.92	0.82	0.86	8144
weighted avg	0.92	0.92	0.91	8144

0.9161345776031434

Decision Tree

```
[44]: from sklearn.tree import DecisionTreeClassifier

DTc = DecisionTreeClassifier()
DTc.fit(X_train_transformed,Y_train)
y_pred_DT = DTc.predict(X_test_transformed)

print(classification_report(Y_test, DTc.predict(X_test_transformed)))
DT_AUC=metrics.accuracy_score(Y_test,y_pred_DT)
print(DT_AUC)
```

	precision	recall	f1-score	support
0	0.94	0.92	0.93	6423
1	0.73	0.78	0.75	1721
accuracy			0.89	8144
macro avg	0.84	0.85	0.84	8144
weighted avg	0.90	0.89	0.89	8144

0.893172888015717

Random Forest

```
[45]: from sklearn.ensemble import RandomForestClassifier

RFC = RandomForestClassifier()
RFC.fit(X_train_transformed,Y_train)
y_pred_RF = RFC.predict(X_test_transformed)

print(classification_report(Y_test, RFC.predict(X_test_transformed)))
RFC_AUC=metrics.accuracy_score(Y_test,y_pred_RF)
print(RFC_AUC)
```

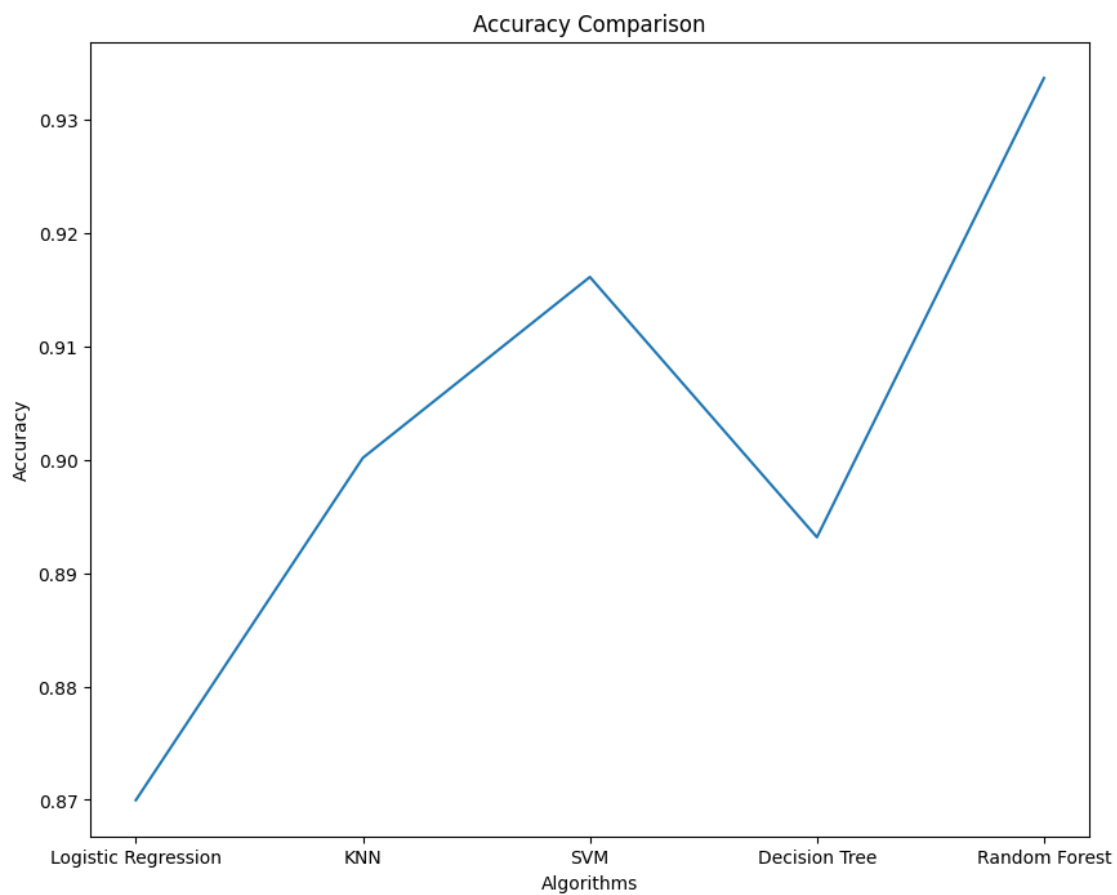
	precision	recall	f1-score	support
0	0.93	0.99	0.96	6423
1	0.95	0.73	0.82	1721

accuracy			0.93	8144
macro avg	0.94	0.86	0.89	8144
weighted avg	0.93	0.93	0.93	8144

0.9336935166994106


Plot of Accuracy of models

```
[46]: plt.figure(figsize=(10,8))
plt.plot(['Logistic Regression','KNN','SVM','Decision Tree','Random_
Forest'], [LR_AUC,KNN_AUC,SVM_AUC,DT_AUC,RFC_AUC])
plt.xlabel('Algorithms')
plt.ylabel('Accuracy')
plt.title('Accuracy Comparison')
plt.show()
```



Bar plot of Accuracy

```
[99]: models = [LR_model, KNNcls, SVMa, DTc, RFC]
models_names = ['Logistic Regression', 'KNN', 'SVM', 'Decision Tree', 'Random_
↳Forest']
accuracies = [LR_AUC, KNN_AUC, SVM_AUC, DT_AUC, RFC_AUC]
plt.bar(models_names, accuracies)
plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.title('Accuracy of Different Models')
plt.show()
```



Conclusion: After evaluating the given models the accuracy of Random Forest Classifier gives an accuracy of 0.935, compared to other models and Support vector machines with 0.916 accuracy.

```
[ ]:
```